

Smart Blood Donation System

A Case Study in Applied Data Structures & Algorithms

Engineering for Millisecond Response in Critical Scenarios

High-Performance Algorithms • Custom Data Structures • Real-World Application

The Mandate: From Hours to Seconds When Lives Are at Stake



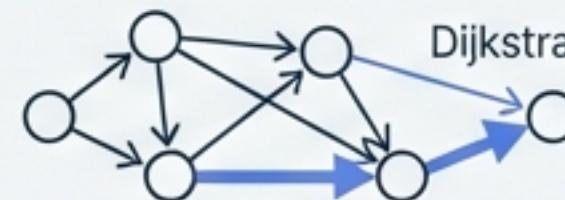
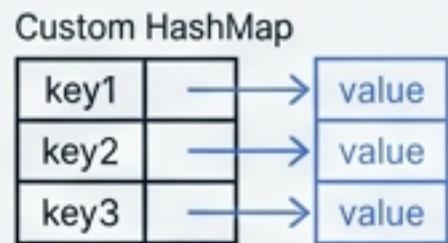
The Problem

- Blood shortages and slow manual search processes lead to **preventable deaths**.
- In emergencies, searching for donors is a race against **time**, where **hours of manual calls** are the norm.

The Engineered Solution

- An automated, **intelligent** system that connects donors with recipients instantly.
- The core mission is to **leverage custom Data Structures and Algorithms** to eliminate search latency.

The Architectural Goals



- **O(1) Donor Lookup:** Instantaneous retrieval of donors by blood type using a custom HashMap.

- **Geographic Optimization:** Find the nearest available donor, not just any donor, using a Graph and Dijkstra's algorithm.

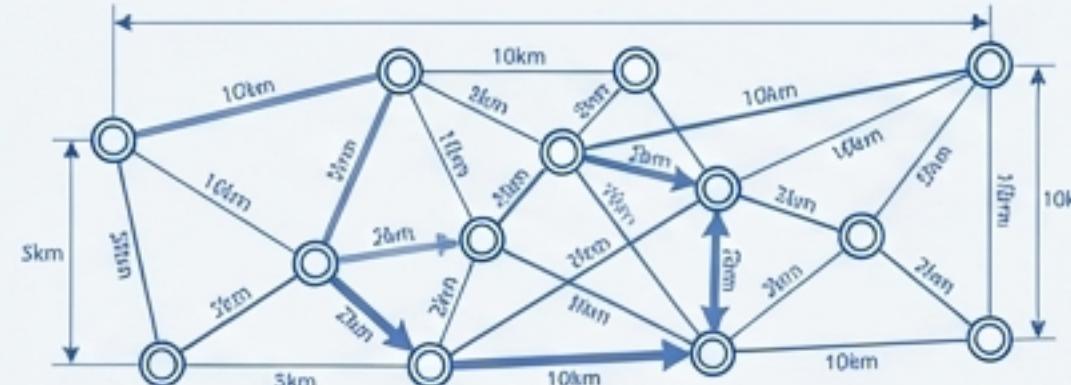
- **Efficient History:** Maintain immutable transaction records with O(1) insertion via a custom LinkedList.

The Core Challenge: Why a Linear Search is Not an Option

With a dataset of over 10,000 users, a linear search ($O(n)$) is dangerously slow in an emergency. Every second of computation is a second lost for a patient.

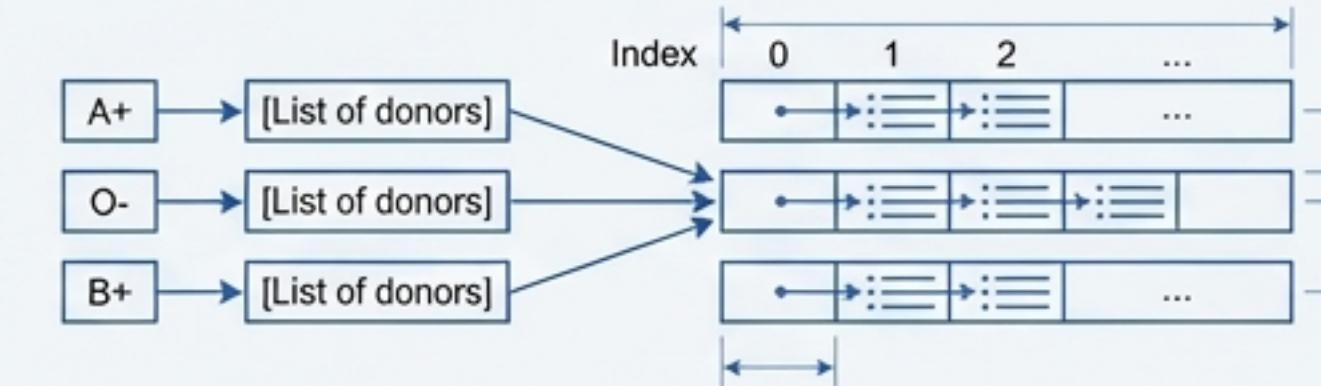
The DSA-Driven Architecture: Our solution is built on a foundation of custom data structures, each engineered for a specific performance objective.

Geographic Matching



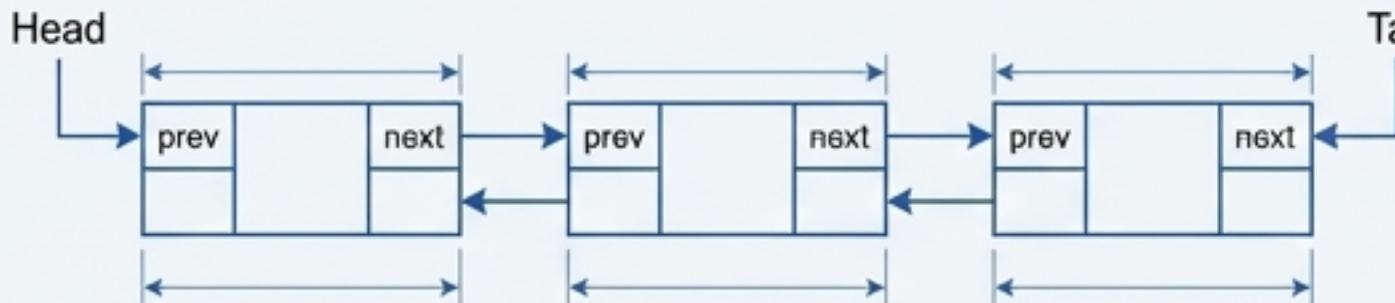
Custom Graph with Dijkstra's algorithm

Blood Type Matching



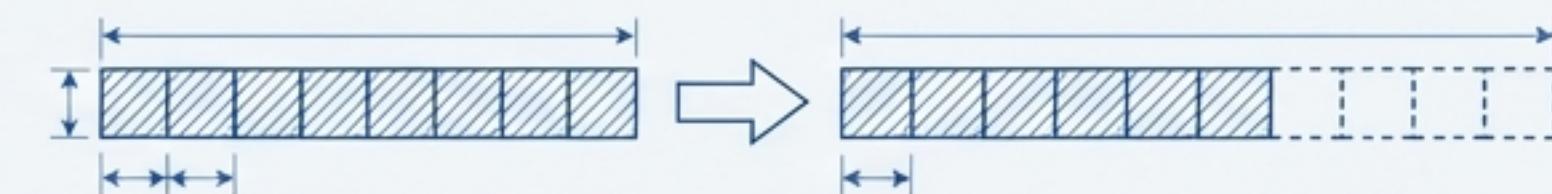
Custom HashMap for ' $O(1)$ ' lookup

History Tracking



Custom LinkedList for ' $O(1)$ ' insertion

Dynamic Storage



Custom Vector to handle variable results

Performance Impact

- Without Custom DSA: ' $O(n)$ ' - Scanning 10,000 donors results in unacceptable delays.
- With Custom DSA: ' $O(1) + O(E \log V)$ ' - Operations complete in milliseconds.

The User-Facing Architecture: Lightweight and Responsive

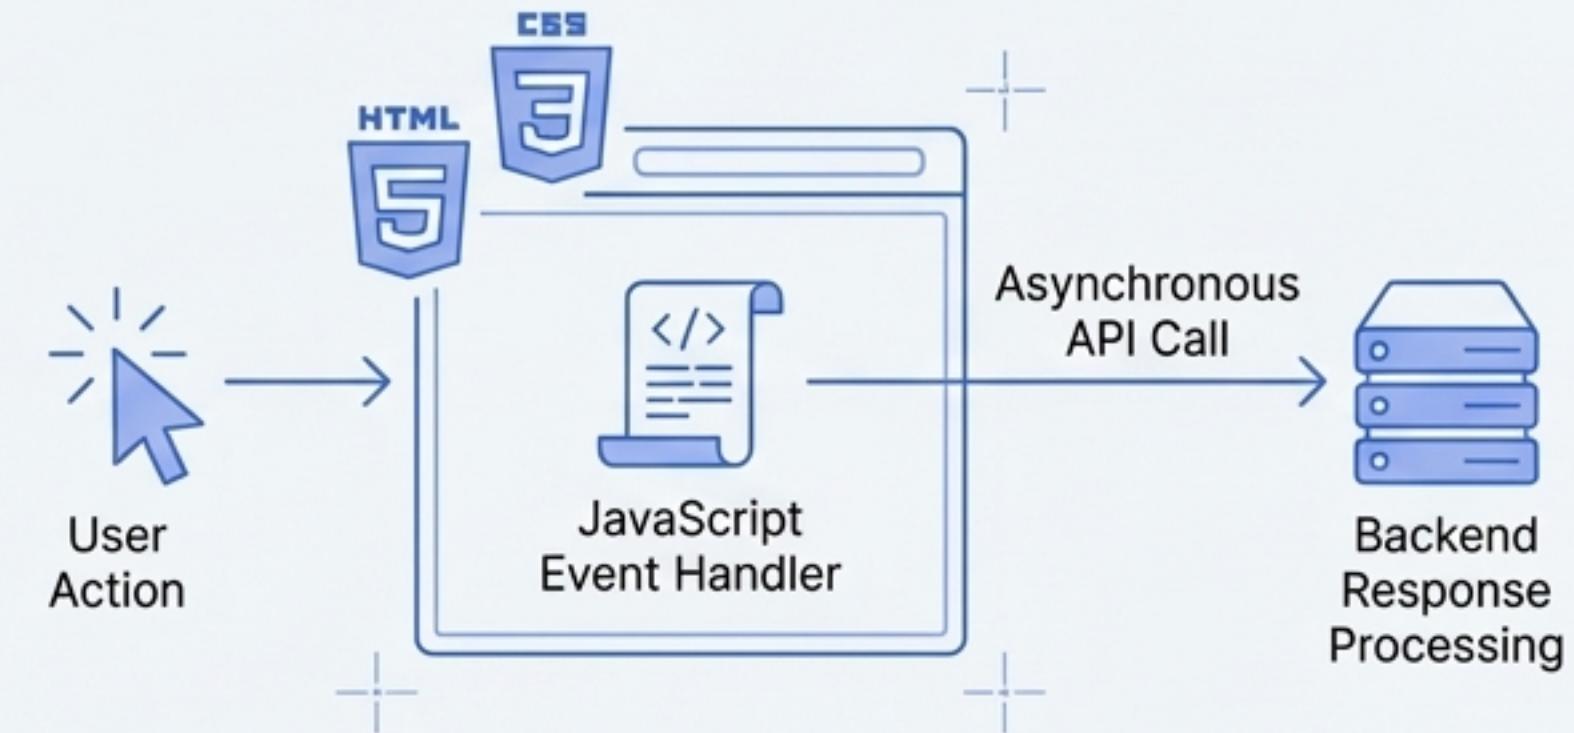
Technology Stack

- **HTML5**: Semantic structure for forms and content.
- **CSS3**: Responsive design, modern gradients, and animations for a clean user experience.
- **Vanilla JavaScript**: Direct DOM manipulation for speed. We intentionally avoided heavy frameworks to minimize load times.

Key Components

- User-facing pages include `index.html`, `login.html`, `register.html`, and `dashboard-donor.html`.
- Client-side session state is managed using `localStorage`.

Interaction Flow

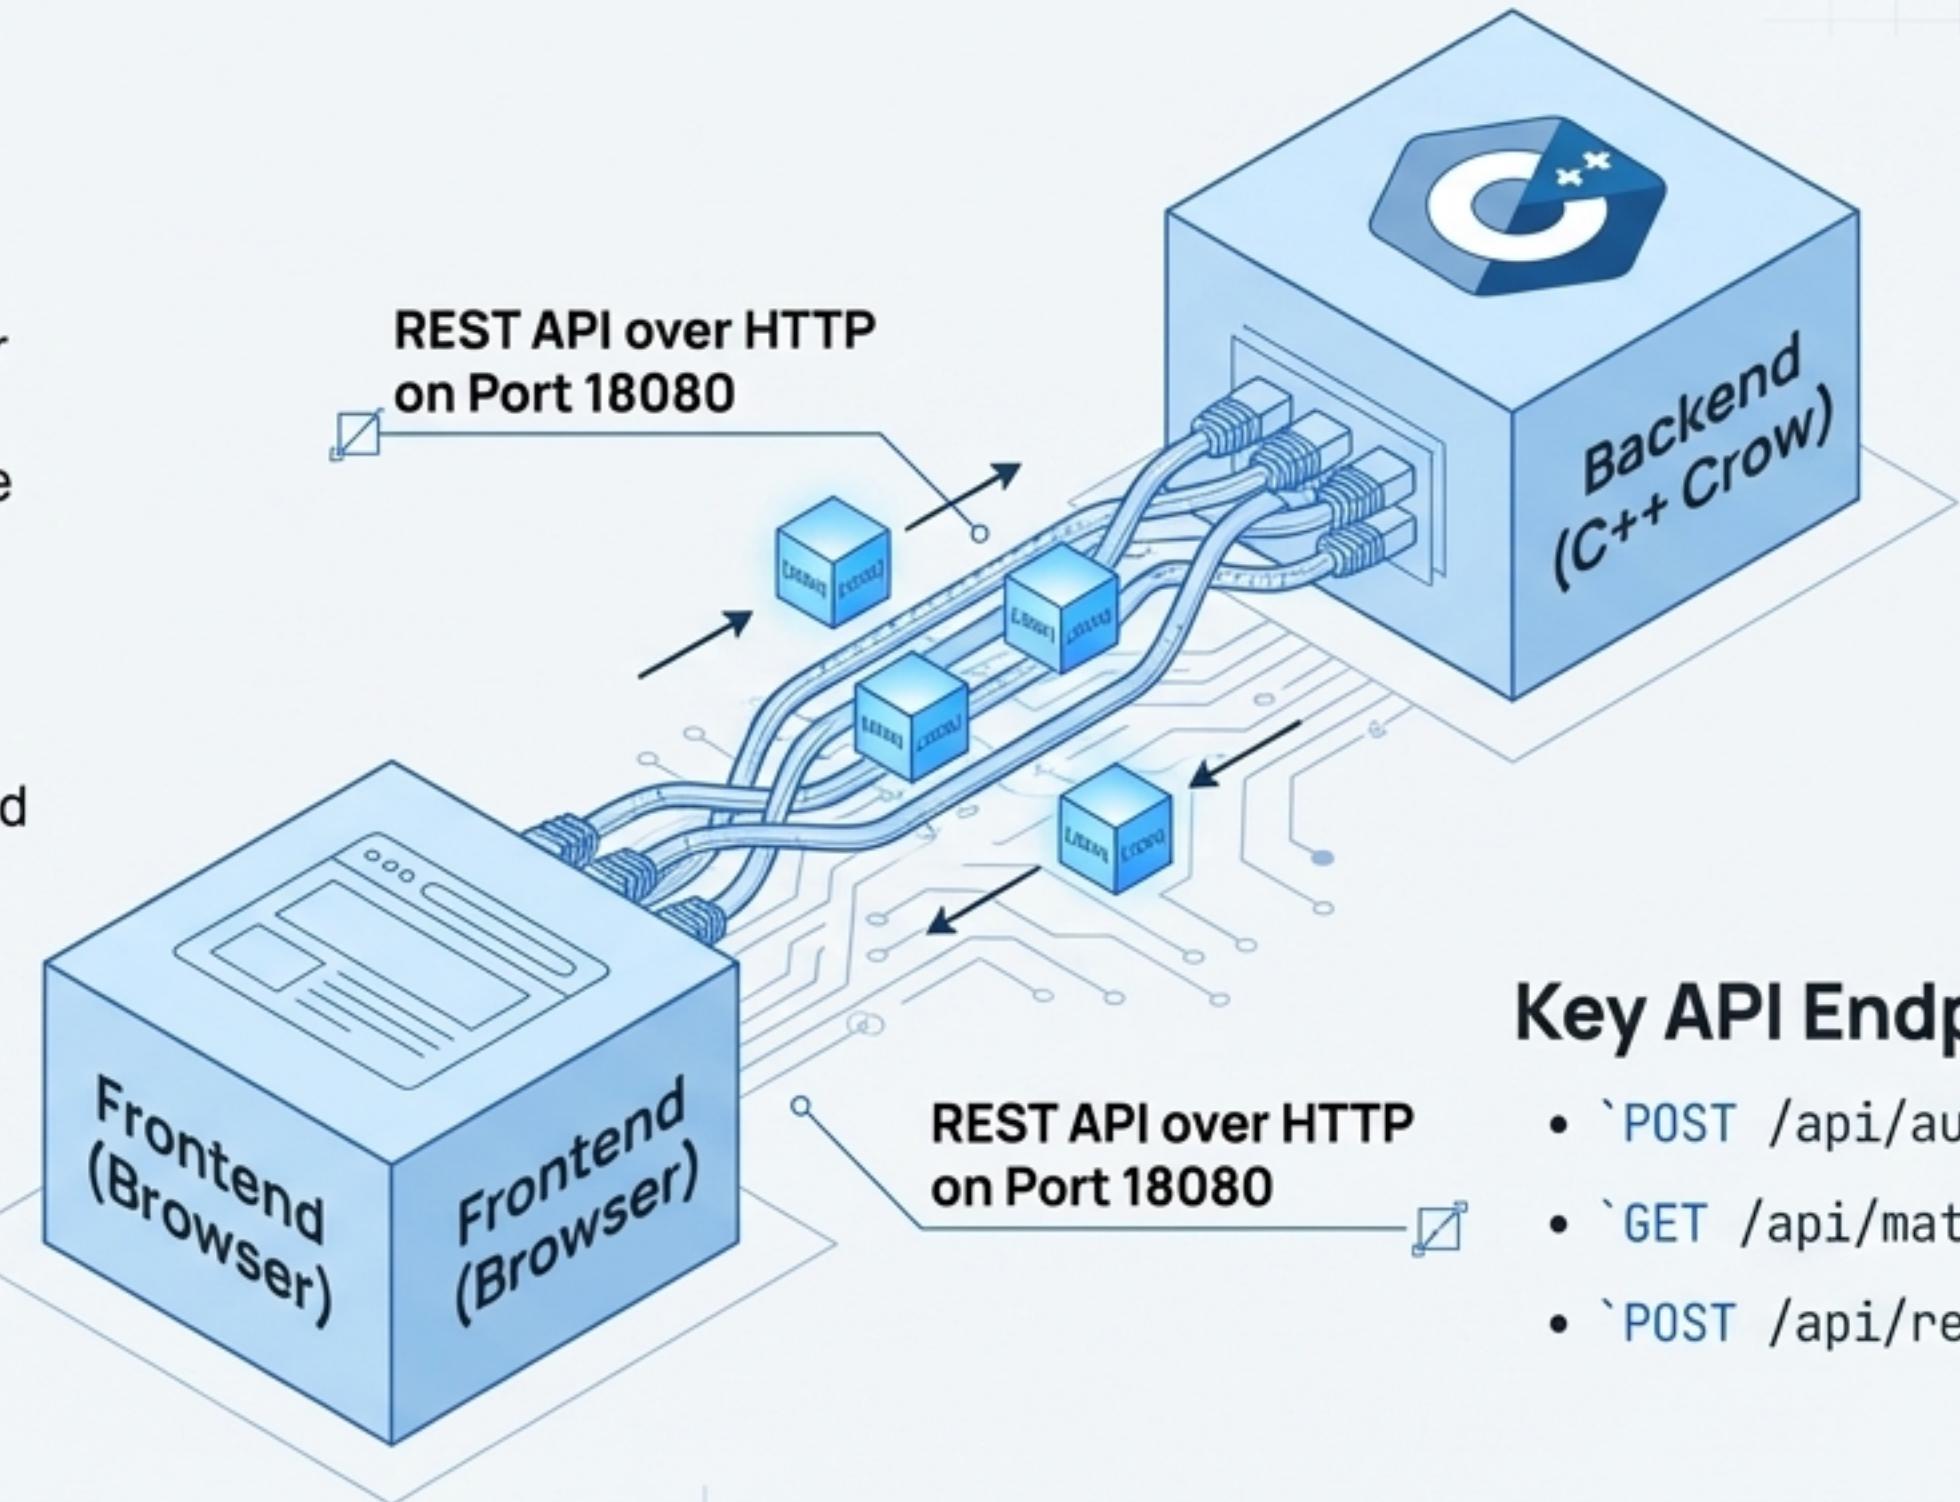


A clear, linear process: User Action -> JavaScript Event Handler -> Asynchronous API Call -> Backend Response Processing.

The Communication Bridge: REST API & C++ Backend

Justification for Technology Choices

- **Why Crow?:** A modern C++17 framework built on ASIO, chosen for its high-performance, low-overhead characteristics, which are essential for our speed requirements.
- **Why REST?:**
 - **Simplicity:** More straightforward than GraphQL for the project's well-defined scope.
 - **Native Support:** Works seamlessly with the browser's native `fetch()` API.
 - **Testability:** Easily tested and debugged using standard tools like Postman and cURL.

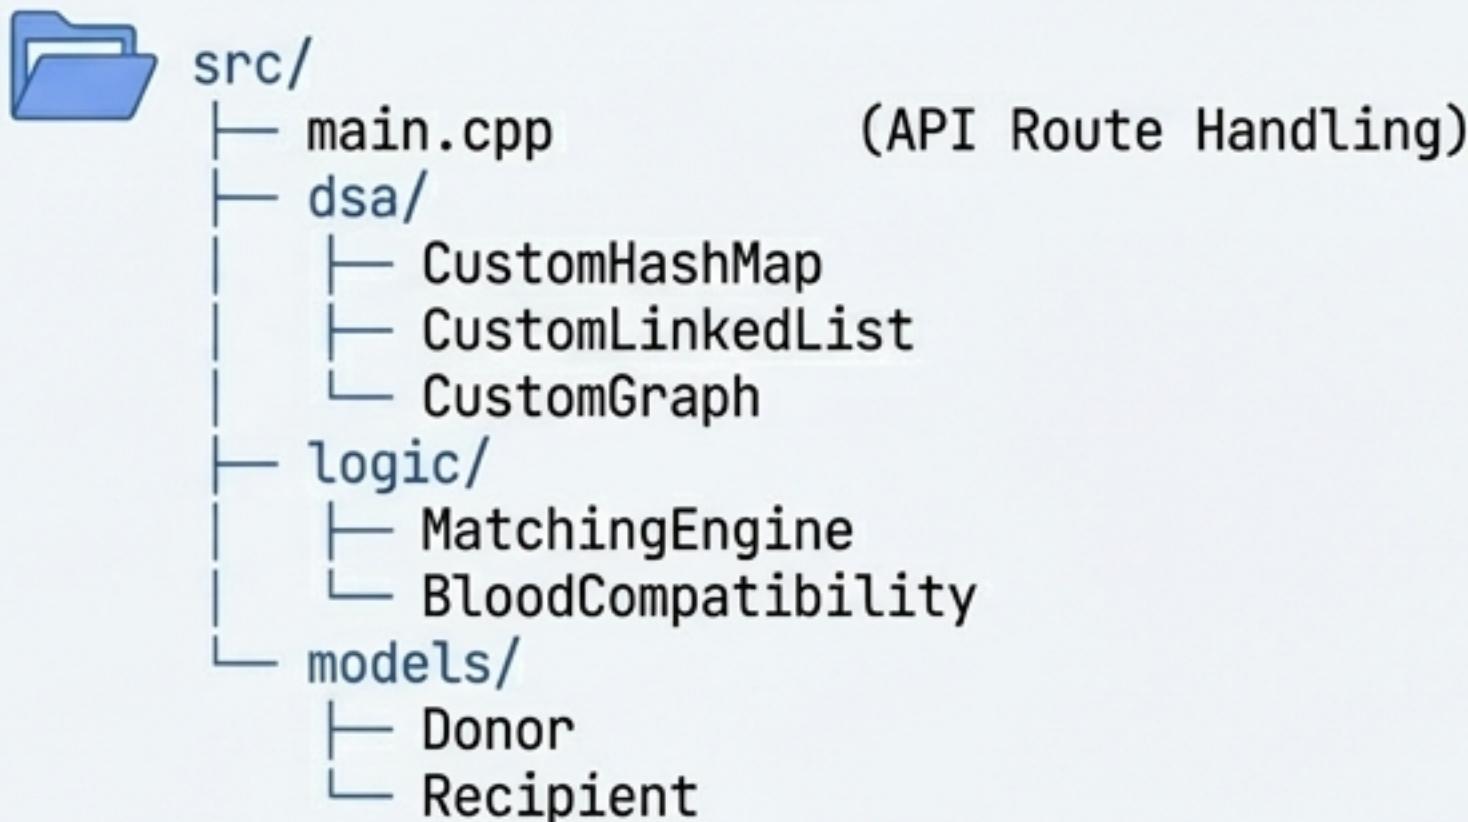


Key API Endpoints

- `POST /api/auth/register`
- `GET /api/match/find`
- `POST /api/request/blood`

The Backend Blueprint: Code Structure and Data Flow

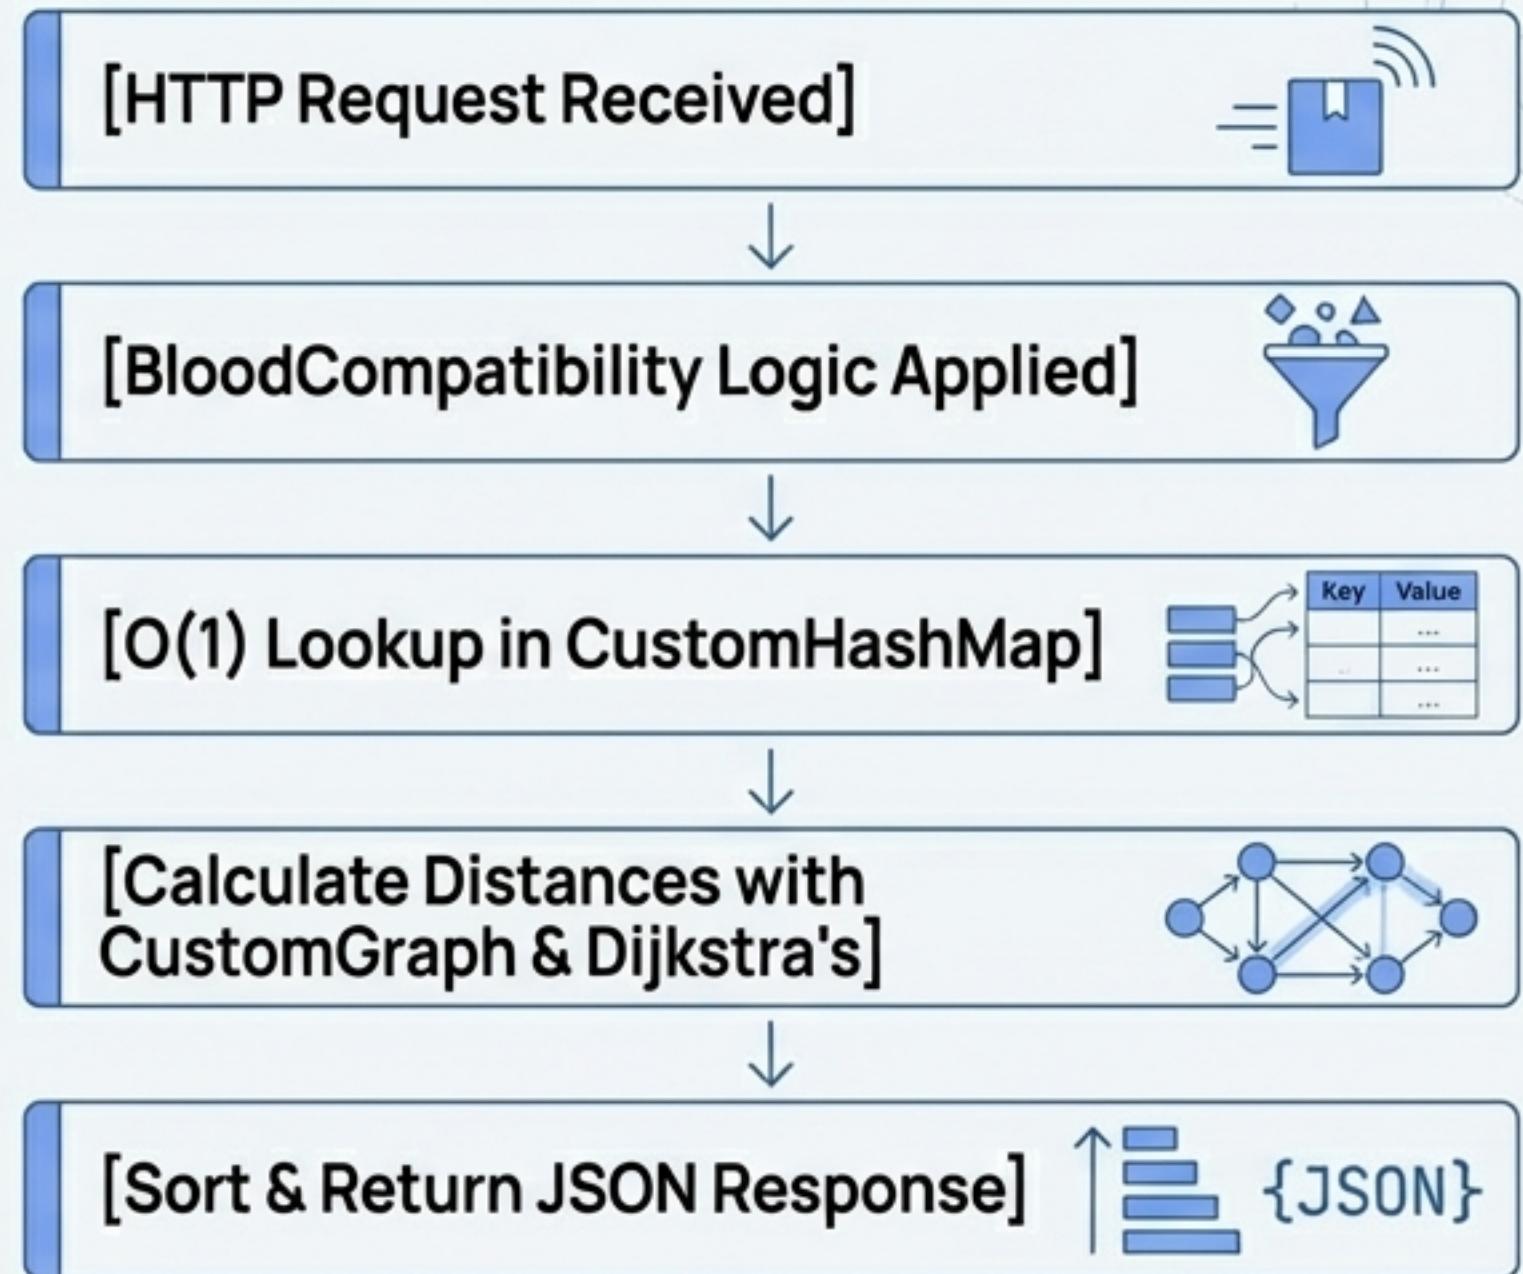
Project Directory Structure



Data Flow Example 1: User Registration

1. HTTP Request hits `main.cpp`.
2. A `Donor` object is created.
3. The object is inserted into the `CustomHashMap` in memory and persisted via a `CSVHandler`.

Data Flow Example 2: Finding a Match



Component Specifications: An Overview of Core Algorithms

Key Algorithms Employed

1. **Hash Function & Linear Probing**: For $O(1)$ average time donor lookup.
2. **Dijkstra's Algorithm**: For $O((V+E)\log V)$ shortest path calculation to find the nearest donor.
3. **Breadth-First Search (BFS)**: For $O(V+E)$ graph traversal.
4. **Doubly Linked List**: For $O(1)$ insertion of transaction history.
5. **Merge Concept**: For $O(n)$ history traversal and reporting.

Performance Comparison Table

Algorithm	Use Case	Time Complexity	Space Complexity
Custom HashMap	Donor Lookup by Email	$O(1)$ Avg, $O(n)$ Worst	$O(n)$
Dijkstra's	Find Nearest Donor	$O(E \log V)$	$O(V)$
Custom LinkedList	Transaction History	$O(1)$ Insertion	$O(n)$

Custom Component #1: The Instant-Access Indexing System (HashMap)

Engineering Purpose

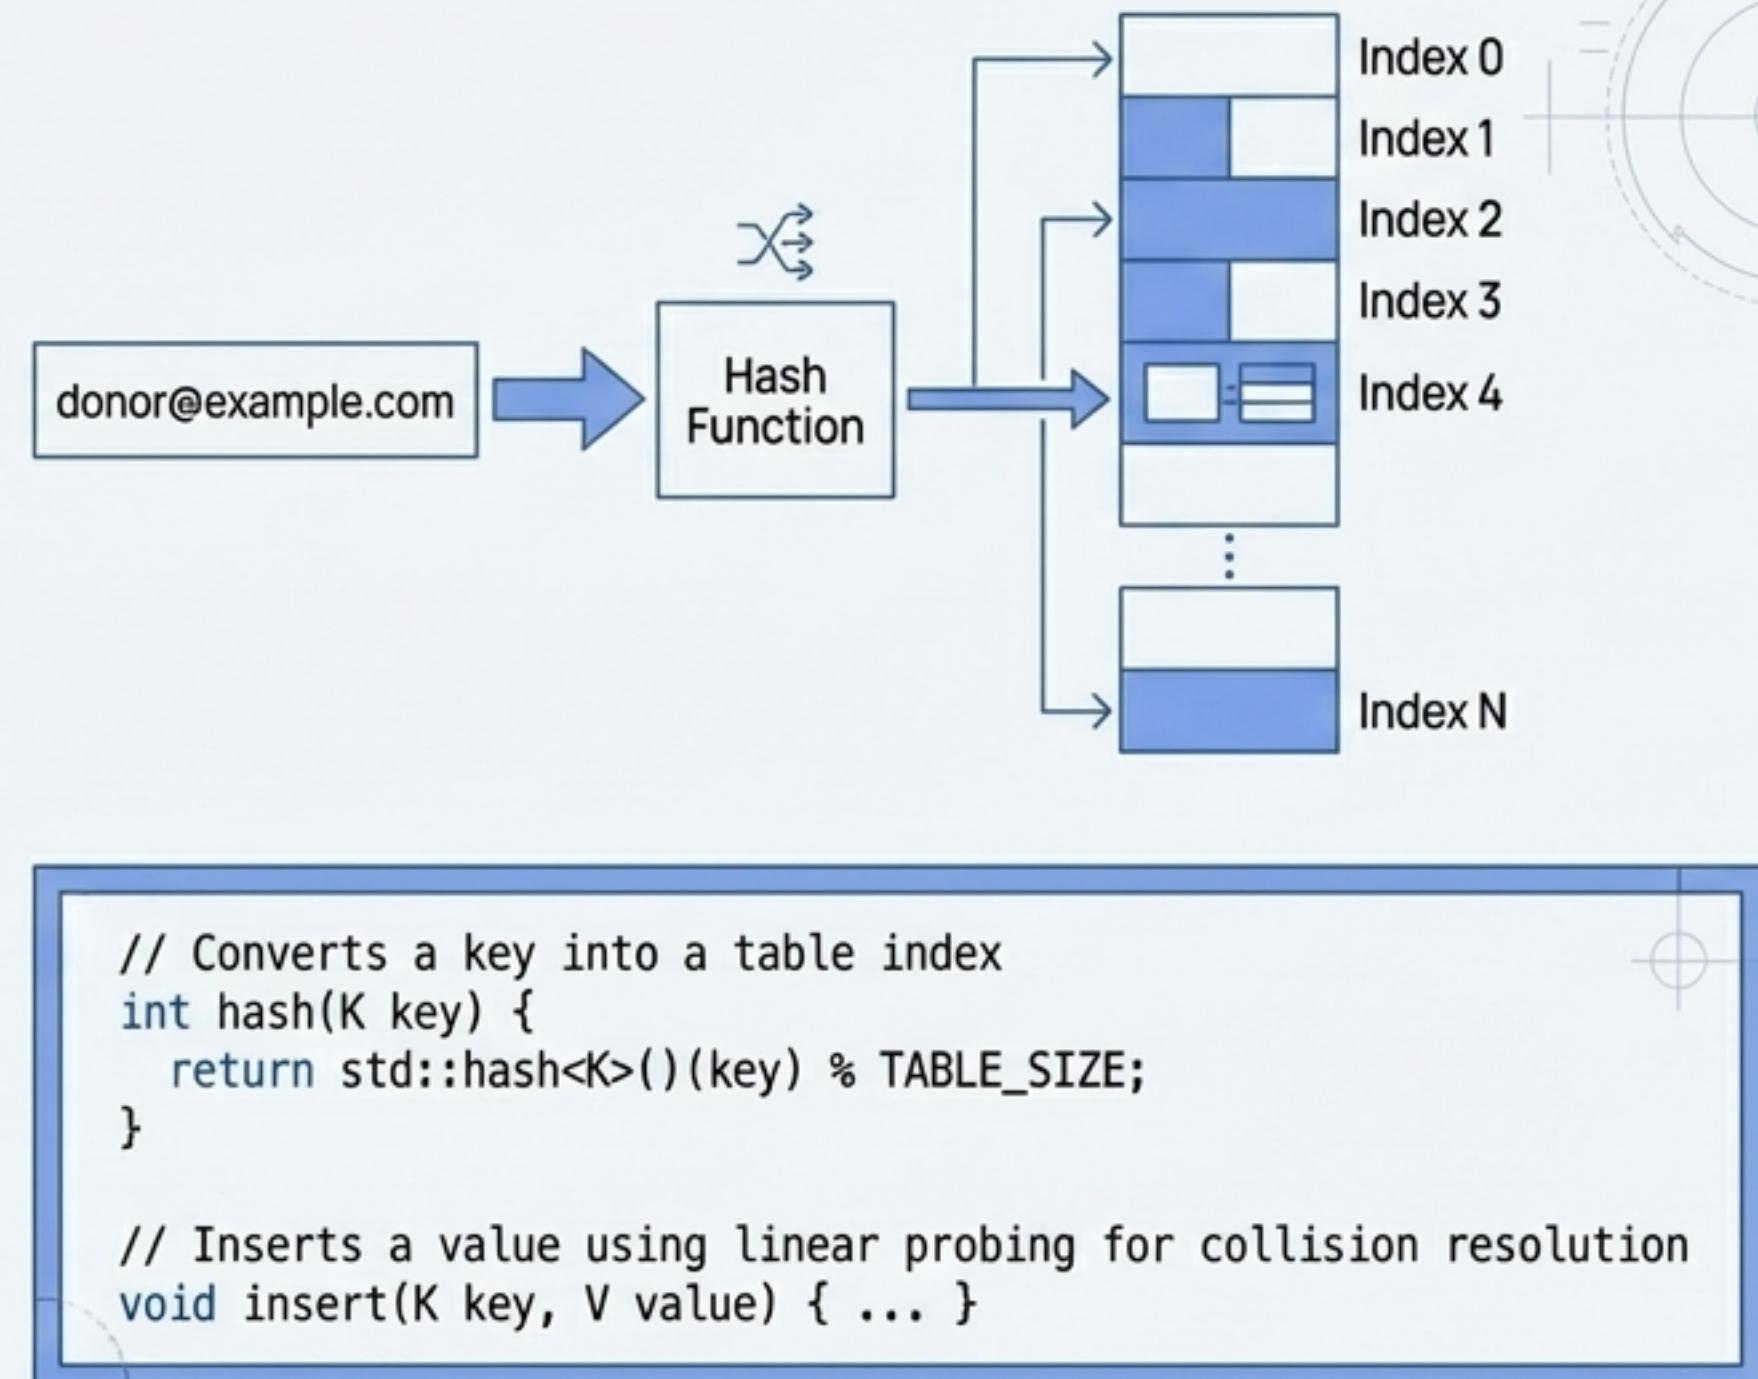
To find any donor by their unique email in constant time, ' $O(1)$ '. This is critical for fast dashboard lookups and authentication.

Technical Implementation

- A hash table implemented with an array of key-value pairs.
- **Collision Handling:** Uses linear probing to resolve hash collisions efficiently.
- **Hash Function:** A custom function converts the string-based email key into an integer array index.

Performance

- **Best/Average Time Complexity:** ' $O(1)$ '
- **Worst-Case Time Complexity:** ' $O(n)$ ' (in the rare case of extreme collisions)



Custom Component #2: The Immutable Transaction Ledger (LinkedList)

Engineering Purpose

To maintain a chronological history of all donation requests and fulfillments with zero performance overhead on insertion.



Technical Implementation

- A Doubly Linked List, where each node contains data and pointers to both the 'previous' and 'next' nodes.
- **push_front Operation****: Adds a new donation record to the head of the list in ' $O(1)$ ' time.
- **pop_front Operation****: Removes an old or invalid record in ' $O(1)$ ' time.

```
template <typename T>
struct Node {
    T data;
    Node* prev;
    Node* next;
};

void push_front(T data) { ... }
```

Why a LinkedList?

This structure avoids the costly process of array shifting or resizing required by a standard vector when inserting elements at the beginning.

Custom Component #3: The Dynamic Results Container (Vector)

Engineering Purpose

To store a list of matching donors when the final count is unknown beforehand. The container must grow dynamically as more matches are found.

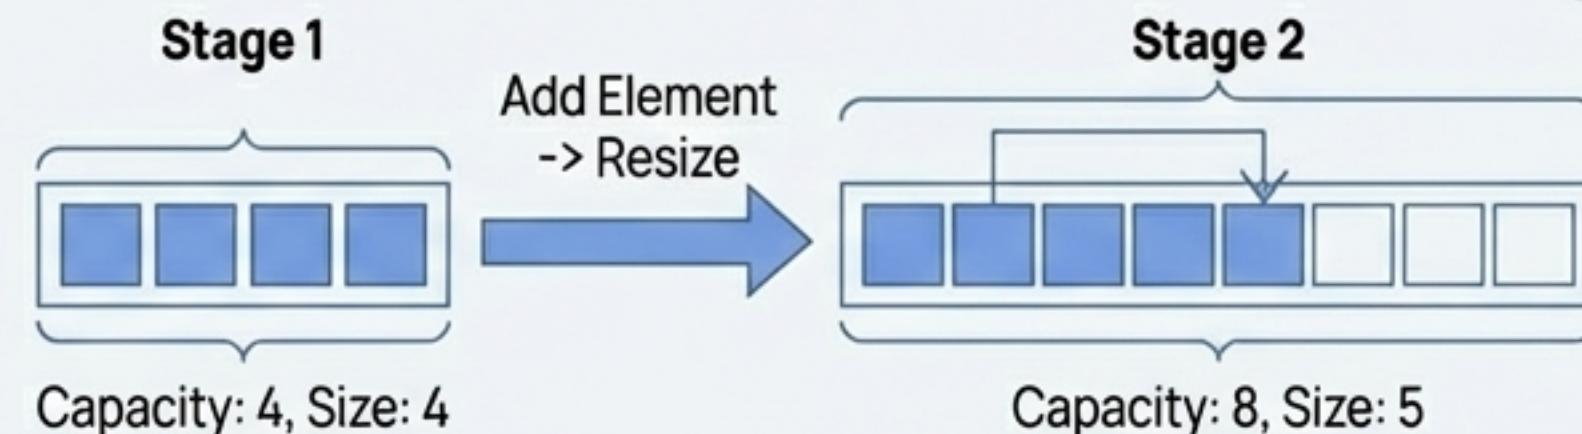
Technical Implementation

- A dynamic array that automatically manages its own storage.
- **Growth Strategy:** Implements an exponential growth model. When the internal array is full, its capacity is doubled (e.g., 10 -> 20 -> 40).

Performance

- **Element Access:** `O(1)`
- **push_back (Adding an element):** `O(1)` amortized time. While resizing is `O(n)`, it happens so infrequently that the average cost per insertion remains constant.

****Why a Custom Vector?***: It provides the perfect balance of fast access and automatic memory management for returning search results to the frontend.



The Intelligent Logistics Engine: Custom Graph + Dijkstra's Algorithm

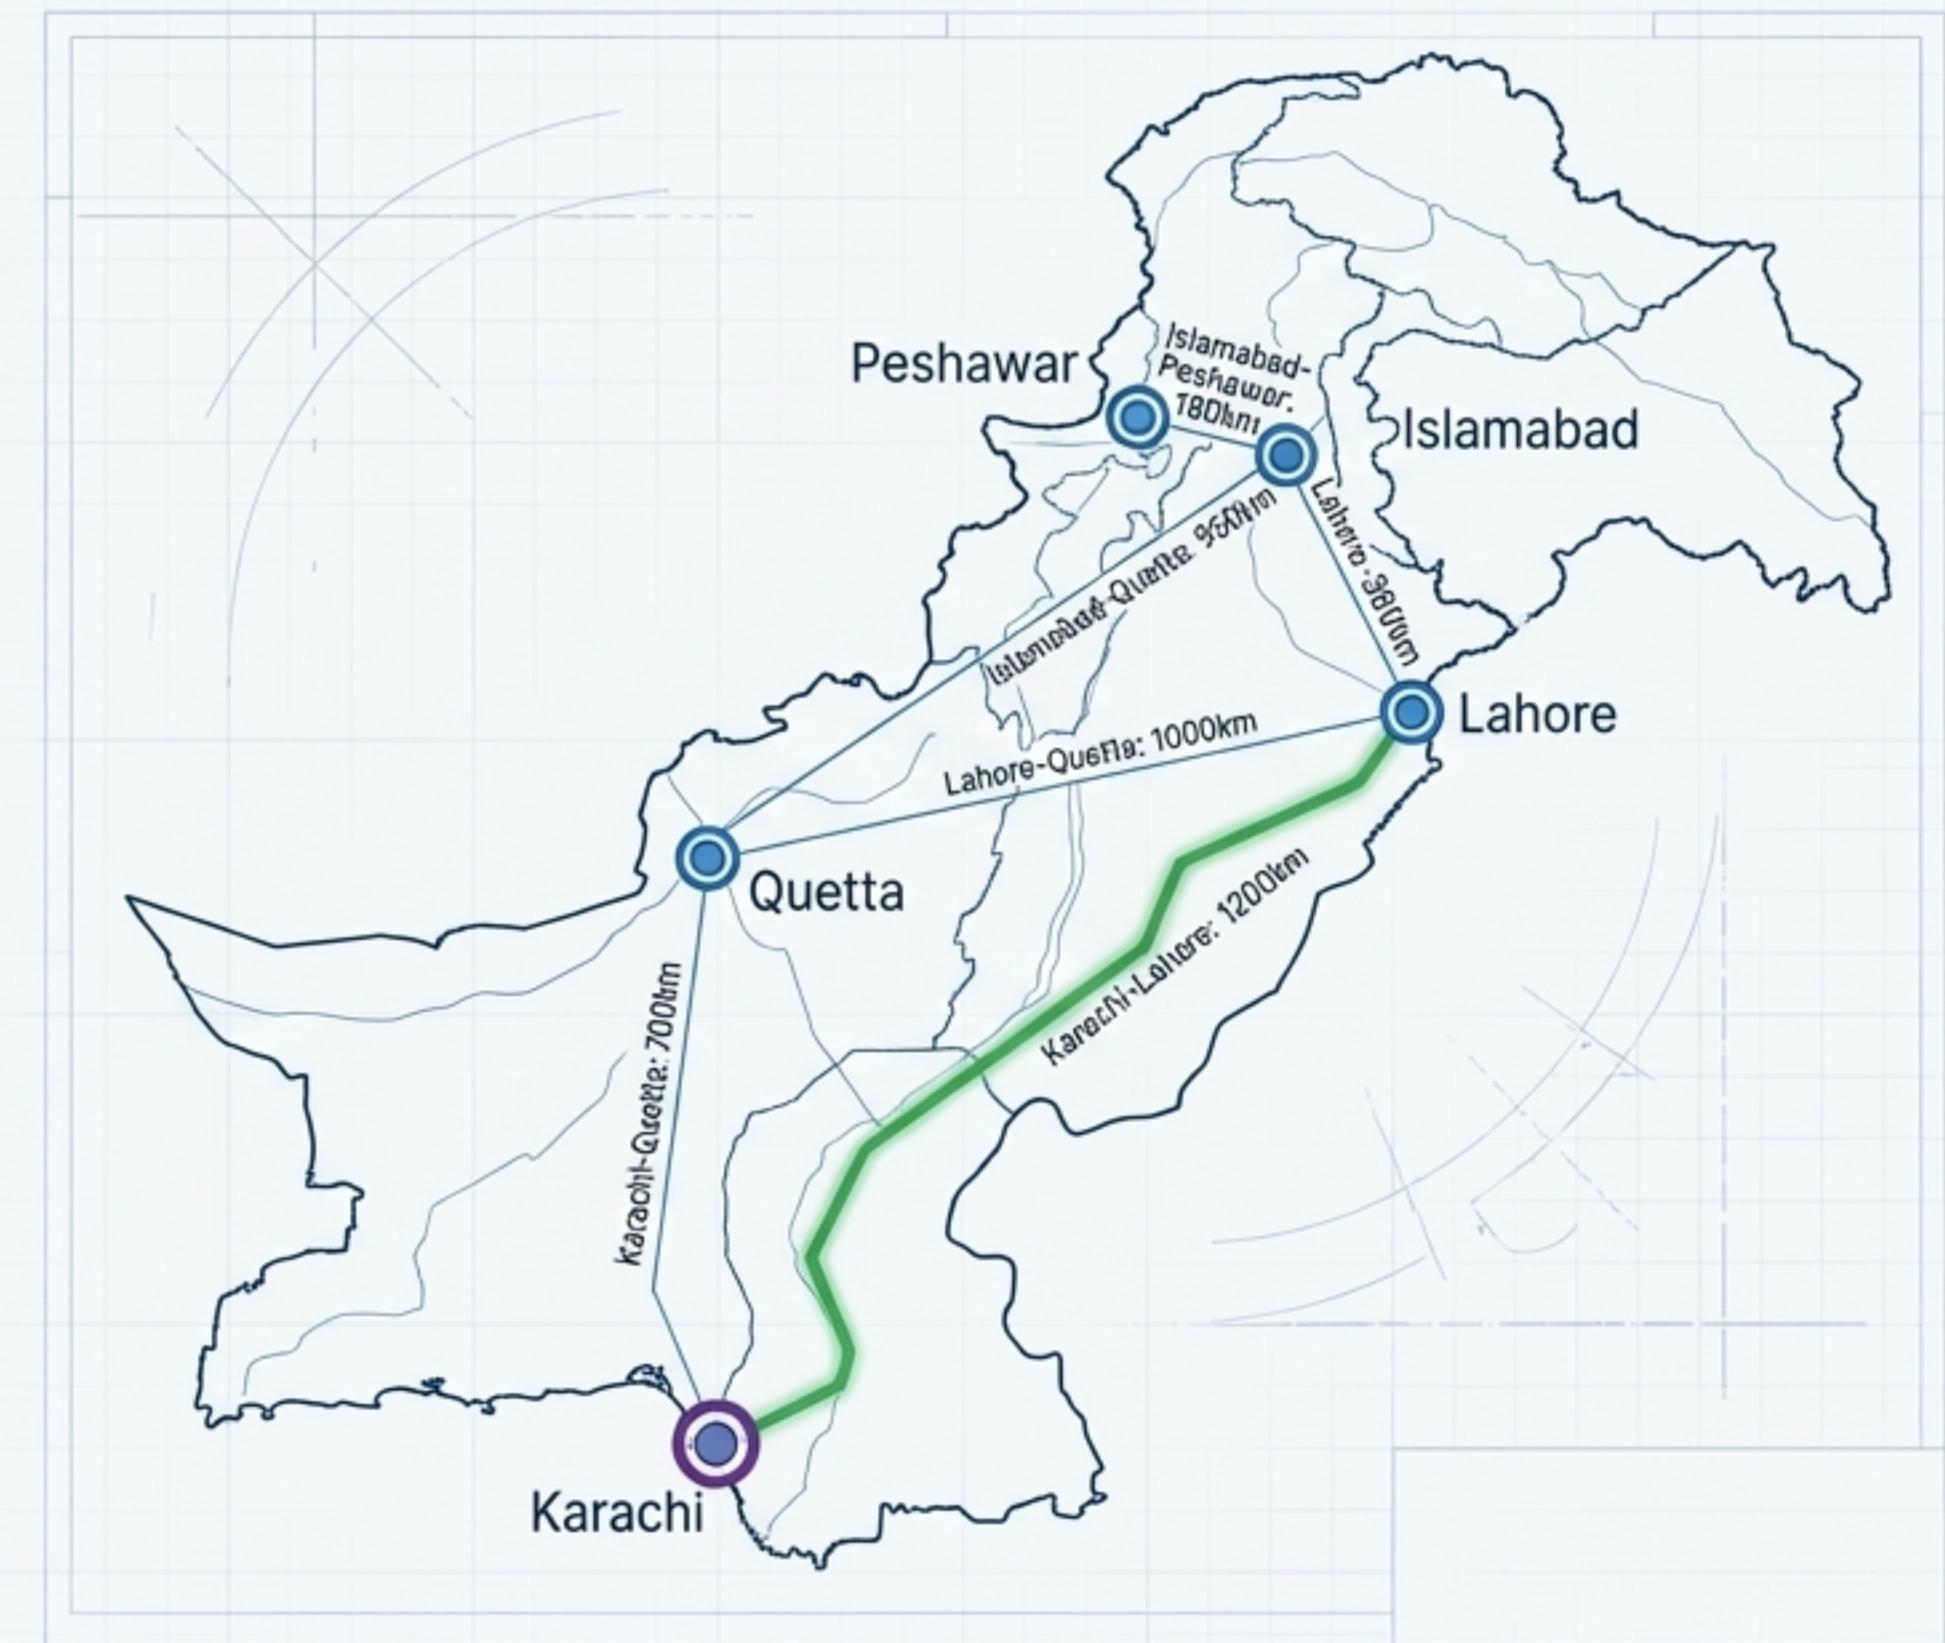
The Core Problem

A patient in Karachi needs O+ blood. We have compatible donors in Lahore, Islamabad, and Quetta. Who is the closest and can respond the fastest?

The Solution: Dijkstra's Shortest Path Algorithm

- 1. Model the Network:** Represent cities as nodes (Vertices) and routes between them as weighted edges (Distances in km).
- 2. Execute Algorithm:** Starting from the patient's city (**the source**), the algorithm greedily explores the graph.
- 3. Priority Queue:** A Min-Heap Priority Queue is used to always select the nearest unvisited city, ensuring optimal pathfinding.
- 4. Update Paths:** Iteratively updates the shortest known distance to all other cities from the source.

****Performance**:** $O(E \log V)$, highly efficient for this type of problem.



The Efficiency Multiplier: Custom Priority Queue (Min-Heap)

****Engineering Purpose****

The Priority Queue is the engine inside Dijkstra's algorithm. Its job is to provide the next closest node to visit in $O(\log n)$ time.

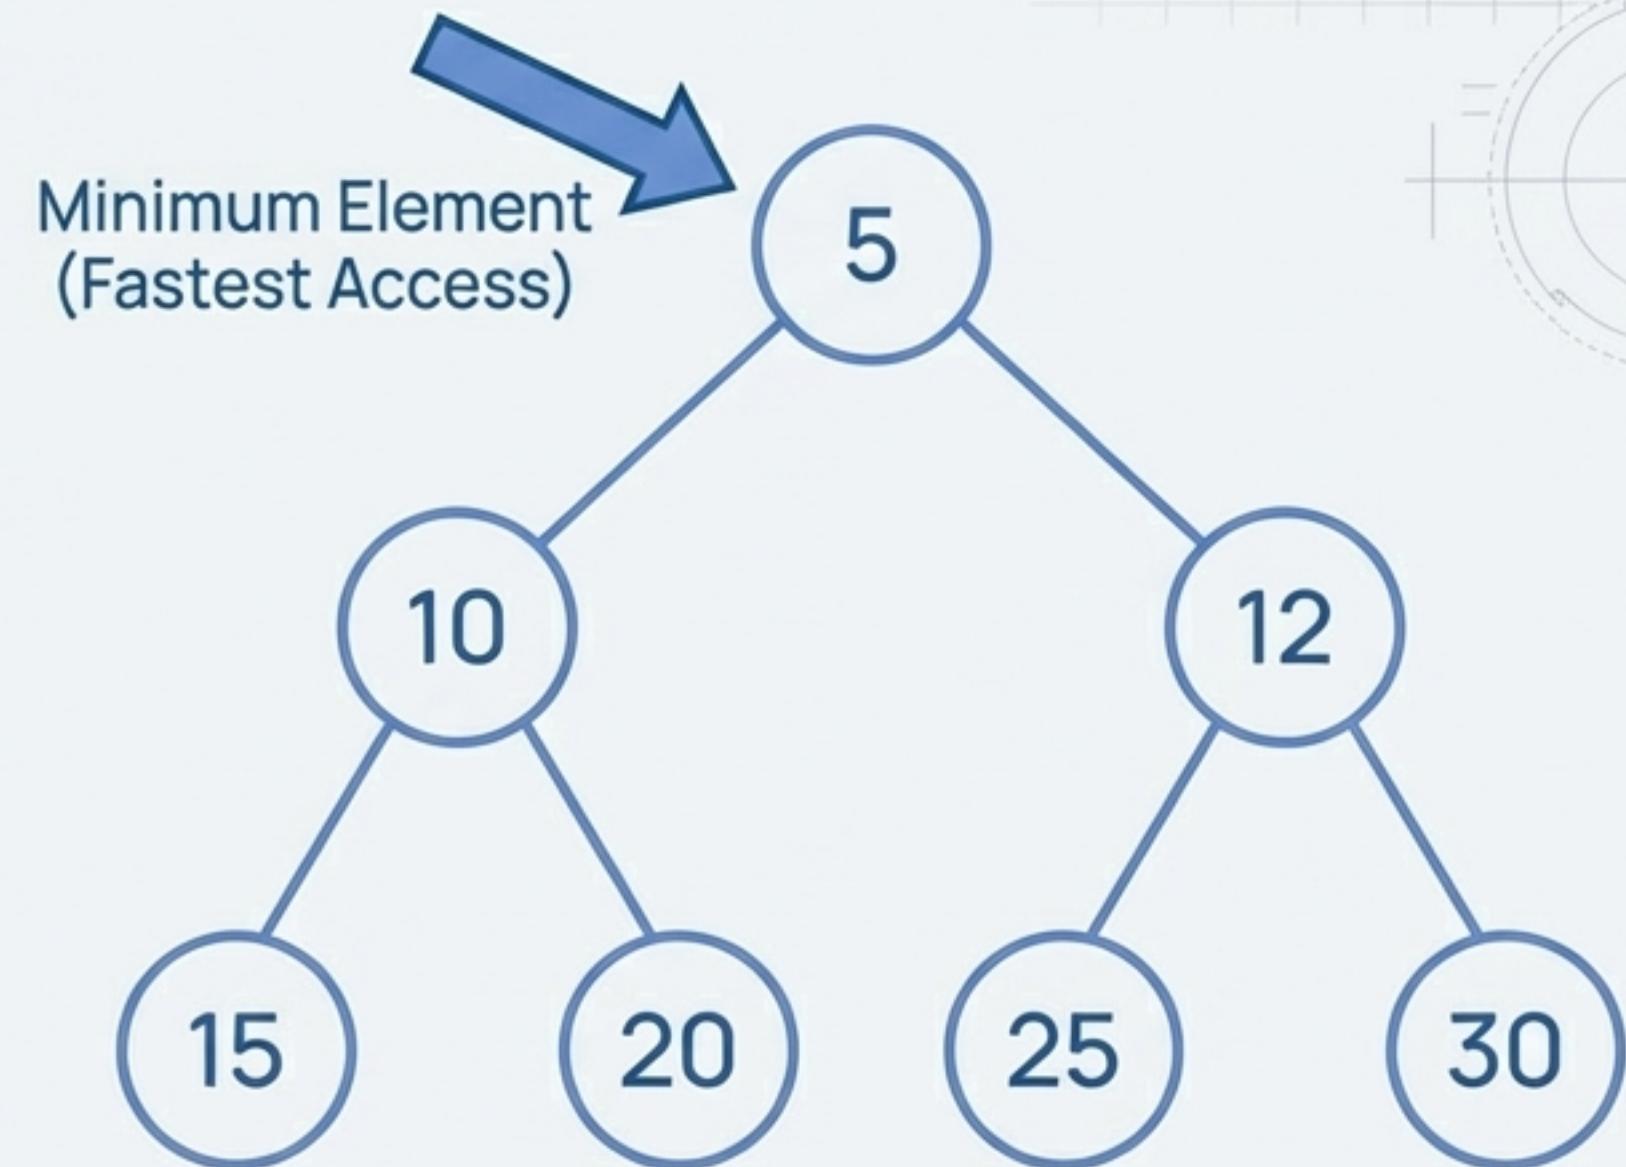
****Technical Implementation****

- A Min-Heap, which is a binary tree with a specific property: a parent node's value is always less than or equal to its children's values.
- This guarantees that the root of the tree is always the element with the minimum value (the shortest distance).

****Performance Impact****

- Push/Pop: $O(\log n)$
- Get Minimum: $O(1)$

****Why is this critical?:** Using a Min-Heap is exponentially faster than scanning an entire list ($O(n)$) at every step of Dijkstra's algorithm to find the minimum distance node. It's a fundamental optimization.



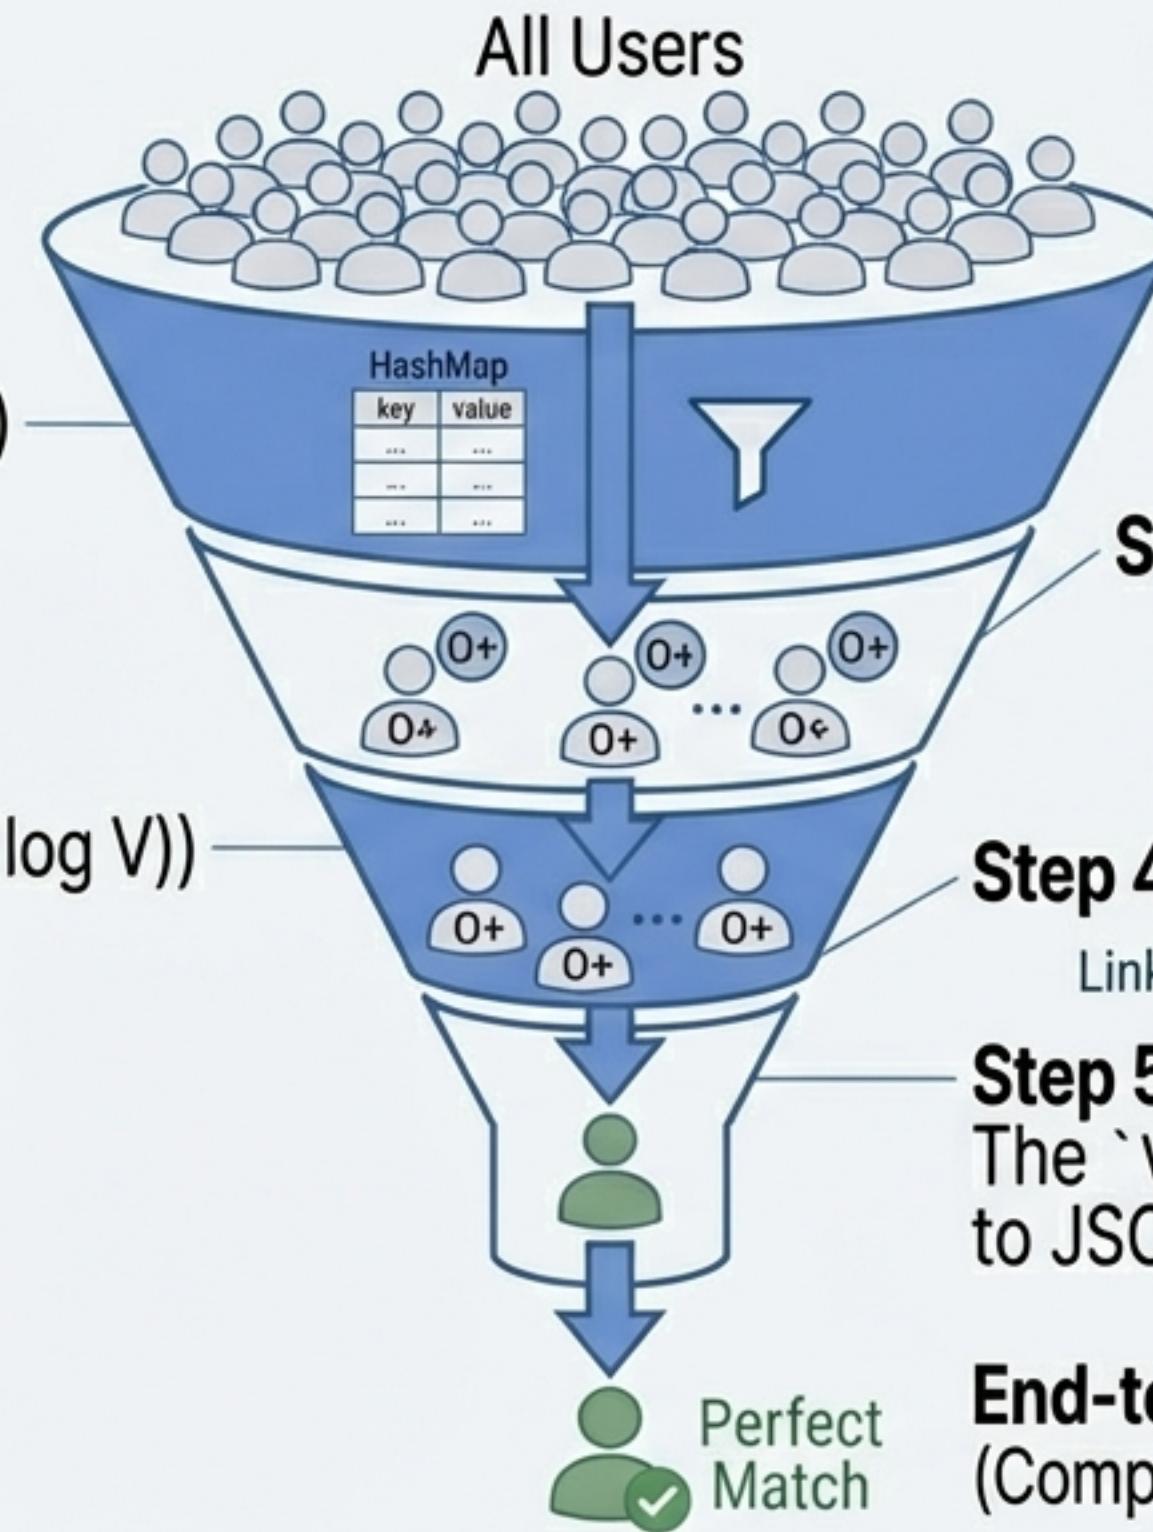
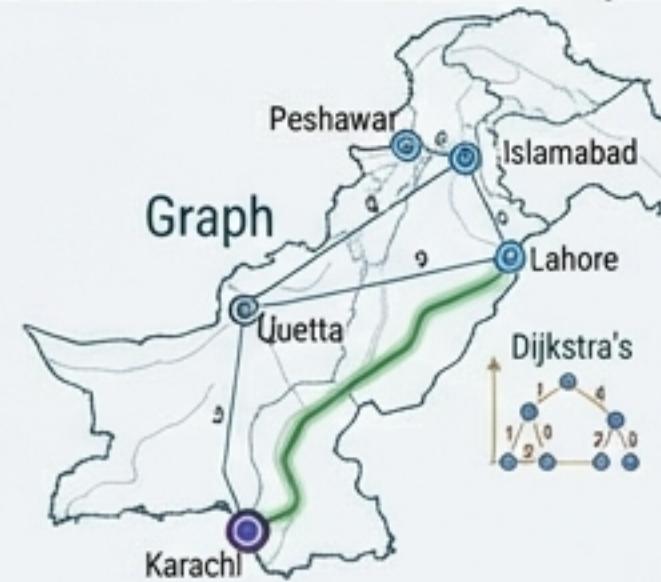
The Integrated System in Motion: A Millisecond Journey

****Scenario**** A patient in Karachi submits an urgent request for O+ blood.

Step 1: Blood Type Filtering ($O(1)$)

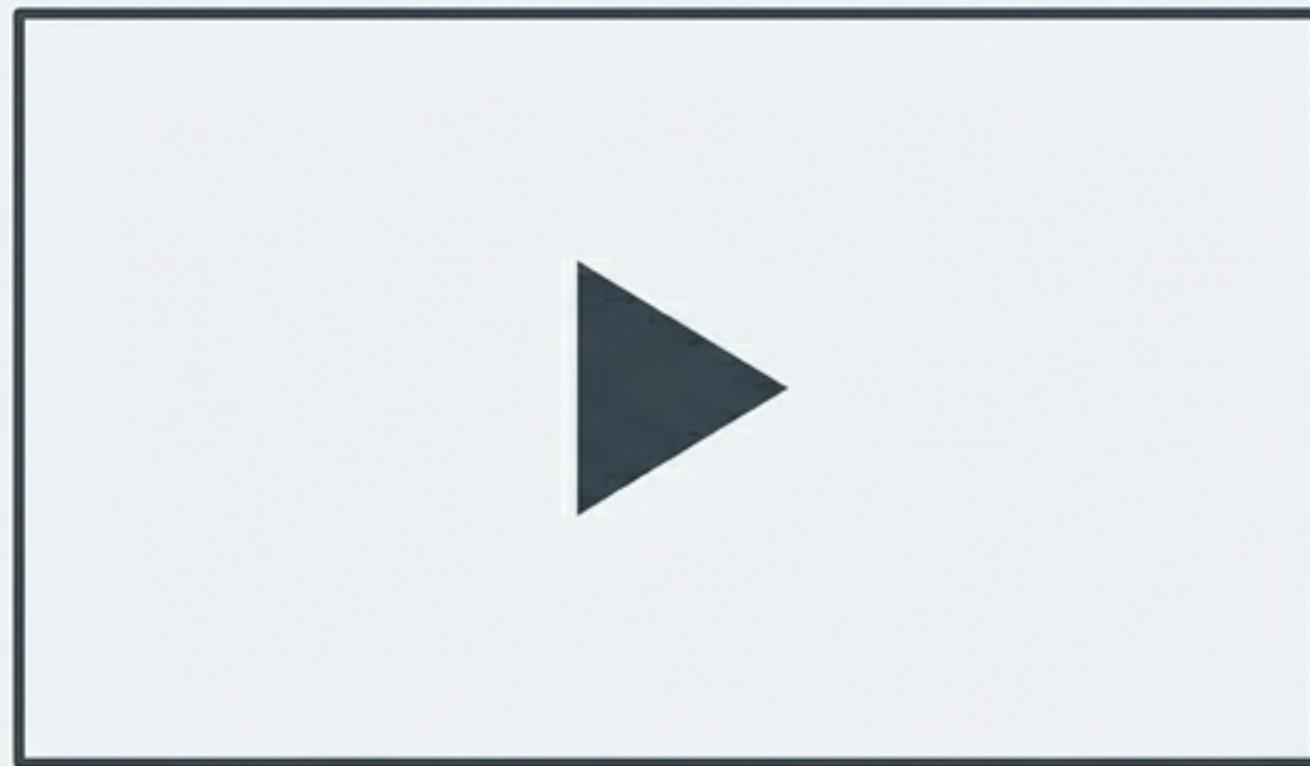
HashMap	
key	value
...	...
...	...

Step 2: Distance Calculation ($O(E \log V)$)



End-to-End Performance: ~2 milliseconds.
(Compared to 2000+ ms for a naive approach).

As-Built Verification: System Demonstration



[VIDEO: DS-Final-Demo.mp4]

****Key Features to Observe in the Demo****

-  **Real-time Matching:** Watch how quickly the system responds to a new blood request.
-  **Geographic Optimization:** Note how the results are correctly sorted by geographic proximity.
-  **Instantaneous UI Updates:** See the seamless performance from user click to result display, proving the efficiency of the full stack.

This is not just theory; it's high-performance algorithms in action, solving a real-world problem.